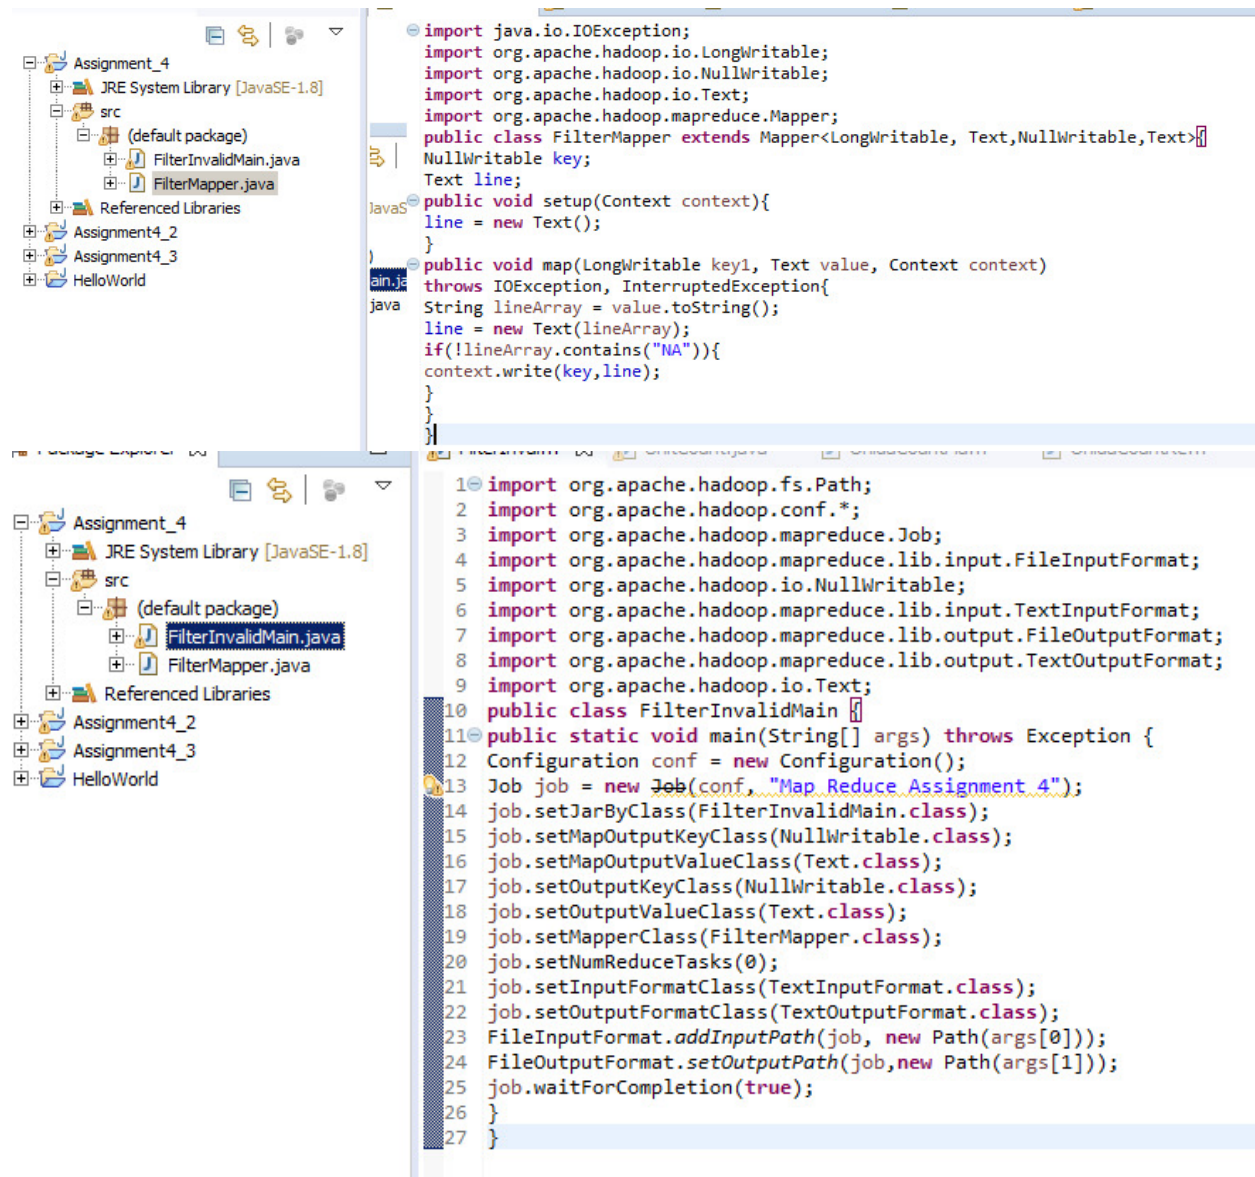


Task 1:

Map Reduce program to filter out Invalid Records; i.e.; records which contain 'NA' in Company or Product name.

Solution:

Code:



```
import java.io.IOException;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class FilterMapper extends Mapper<LongWritable, Text, NullWritable, Text> {
    NullWritable key;
    Text line;

    public void setup(Context context) {
        line = new Text();
    }

    public void map(LongWritable key1, Text value, Context context)
        throws IOException, InterruptedException {
        String lineArray = value.toString();
        line = new Text(lineArray);
        if (!lineArray.contains("NA")) {
            context.write(key, line);
        }
    }
}
```

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.io.Text;

public class FilterInvalidMain {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = new Job(conf, "Map Reduce Assignment 4");
        job.setJarByClass(FilterInvalidMain.class);
        job.setMapOutputKeyClass(NullWritable.class);
        job.setMapOutputValueClass(Text.class);
        job.setOutputKeyClass(NullWritable.class);
        job.setOutputValueClass(Text.class);
        job.setMapperClass(FilterMapper.class);
        job.setNumReduceTasks(0);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.waitForCompletion(true);
    }
}
```

Execution of Jar File:

```

[acadgild@localhost ~]$ hadoop jar Desktop/Jyoti/Assignment4 Task1.jar FilterInvalidMain /television.txt /output4_1
18/07/25 21:32:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
pplicable
18/07/25 21:32:05 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/25 21:32:07 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool in
xecute your application with ToolRunner to remedy this.
18/07/25 21:32:08 INFO input.FileInputFormat: Total input paths to process : 1
18/07/25 21:32:08 INFO mapreduce.JobSubmitter: number of splits:1
18/07/25 21:32:09 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532534151354_0001
18/07/25 21:32:10 INFO impl.YarnClientImpl: Submitted application application_1532534151354_0001
18/07/25 21:32:10 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532534151354_0001/
18/07/25 21:32:10 INFO mapreduce.Job: Running job: job_1532534151354_0001
18/07/25 21:32:39 INFO mapreduce.Job: Job job_1532534151354_0001 running in uber mode : false
18/07/25 21:32:39 INFO mapreduce.Job: map 0% reduce 0%
18/07/25 21:32:51 INFO mapreduce.Job: map 100% reduce 0%
18/07/25 21:32:53 INFO mapreduce.Job: Job job_1532534151354_0001 completed successfully
18/07/25 21:32:53 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=107692
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=834
    HDFS: Number of bytes written=646
    HDFS: Number of read operations=5
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=10175
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=10175
    Total vcore-milliseconds taken by all map tasks=10175
    Total megabyte-milliseconds taken by all map tasks=10419200
  Map-Reduce Framework
    Map input records=18
    Total megabyte-milliseconds taken by all map tasks=10419200
  Map-Reduce Framework
    Map input records=18
    Map output records=16
    Input split bytes=101
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=85
    CPU time spent (ms)=780
    Physical memory (bytes) snapshot=96460800
    Virtual memory (bytes) snapshot=2056757248
    Total committed heap usage (bytes)=32571392
  File Input Format Counters
    Bytes Read=733
  File Output Format Counters
    Bytes Written=646
You have new mail in /var/spool/mail/acadgild

```

Output1:

Output doesn't have records with Product/Company name as 'NA'.

```

[acadgild@localhost ~]$ hadoop fs -ls /output4_1
18/07/25 21:33:16 WARN util.NativeCodeLoader: Unable to load native-
plicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-07-25 21:32 /outp
-rw-r--r-- 1 acadgild supergroup 646 2018-07-25 21:32 /outp
[acadgild@localhost ~]$ hadoop fs -cat /output4_1/part-m-00000
18/07/25 21:33:40 WARN util.NativeCodeLoader: Unable to load native-
plicable
Samsung|Optima|14|Madhya Pradesh|132401|14200
Onida|Lucid|18|Uttar Pradesh|232401|16200
Akai|Decent|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Zen|Super|14|Maharashtra|619082|9200
Samsung|Optima|14|Madhya Pradesh|132401|14200
Onida|Lucid|18|Uttar Pradesh|232401|16200
Onida|Decent|14|Uttar Pradesh|232401|16200
Lava|Attention|20|Assam|454601|24200
Zen|Super|14|Maharashtra|619082|9200
Samsung|Optima|14|Madhya Pradesh|132401|14200
Samsung|Decent|16|Kerala|922401|12200
Lava|Attention|20|Assam|454601|24200
Samsung|Super|14|Maharashtra|619082|9200
Samsung|Super|14|Maharashtra|619082|9200
Samsung|Super|14|Maharashtra|619082|9200
[acadgild@localhost ~]$ hadoop fs -cat /television.txt

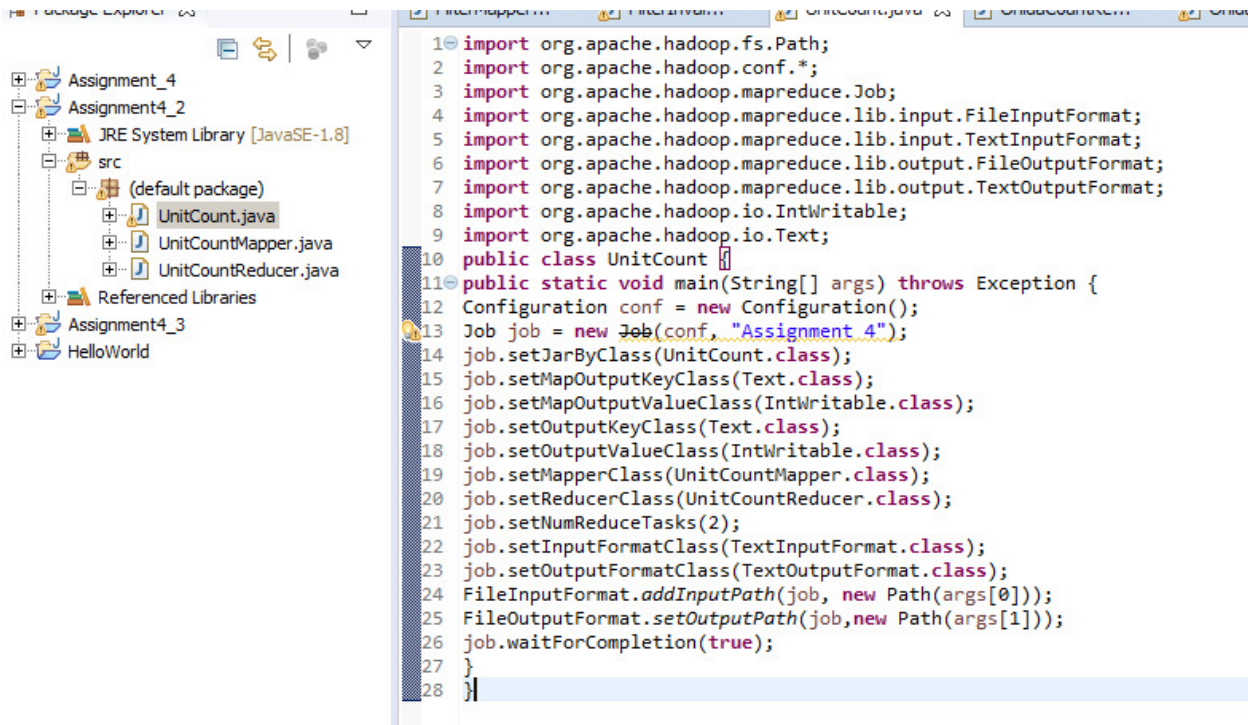
```

Task 2:

Map Reduce program to calculate total units sold for each Company.

Solution:

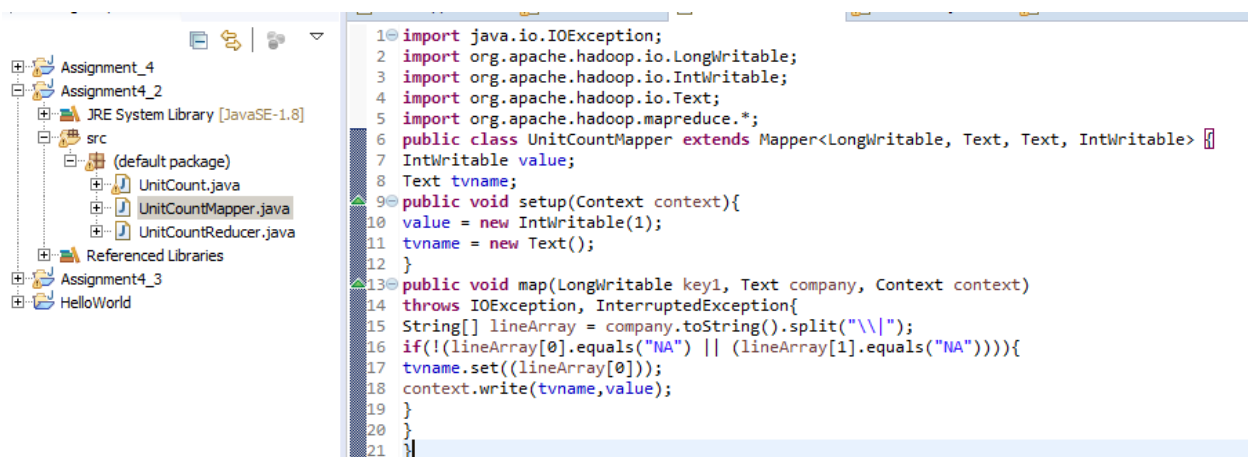
Code:



```

1 import org.apache.hadoop.fs.Path;
2 import org.apache.hadoop.conf.*;
3 import org.apache.hadoop.mapreduce.Job;
4 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
5 import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
6 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
7 import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 public class UnitCount {
11 public static void main(String[] args) throws Exception {
12 Configuration conf = new Configuration();
13 Job job = new Job(conf, "Assignment 4");
14 job.setJarByClass(UnitCount.class);
15 job.setMapOutputKeyClass(Text.class);
16 job.setMapOutputValueClass(IntWritable.class);
17 job.setOutputKeyClass(Text.class);
18 job.setOutputValueClass(IntWritable.class);
19 job.setMapperClass(UnitCountMapper.class);
20 job.setReducerClass(UnitCountReducer.class);
21 job.setNumReduceTasks(2);
22 job.setInputFormatClass(TextInputFormat.class);
23 job.setOutputFormatClass(TextOutputFormat.class);
24 FileInputFormat.addInputPath(job, new Path(args[0]));
25 FileOutputFormat.setOutputPath(job, new Path(args[1]));
26 job.waitForCompletion(true);
27 }
28 }

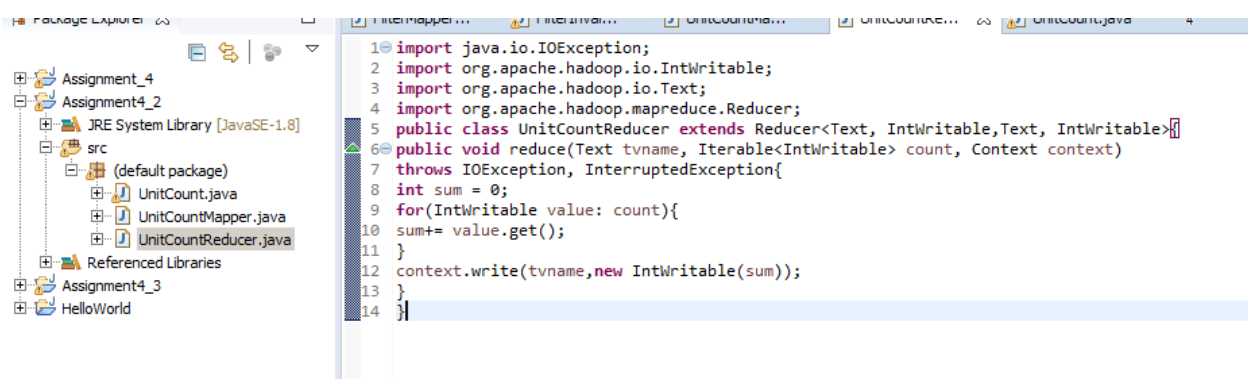
```



```

1 import java.io.IOException;
2 import org.apache.hadoop.io.LongWritable;
3 import org.apache.hadoop.io.IntWritable;
4 import org.apache.hadoop.io.Text;
5 import org.apache.hadoop.mapreduce.*;
6 public class UnitCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
7 IntWritable value;
8 Text tvname;
9 public void setup(Context context){
10 value = new IntWritable(1);
11 tvname = new Text();
12 }
13 public void map(LongWritable key1, Text company, Context context)
14 throws IOException, InterruptedException{
15 String[] lineArray = company.toString().split("\\|");
16 if(!(lineArray[0].equals("NA") || (lineArray[1].equals("NA")))){
17 tvname.set((lineArray[0]));
18 context.write(tvname,value);
19 }
20 }
21 }

```



```

1 import java.io.IOException;
2 import org.apache.hadoop.io.IntWritable;
3 import org.apache.hadoop.io.Text;
4 import org.apache.hadoop.mapreduce.Reducer;
5 public class UnitCountReducer extends Reducer<Text, IntWritable,Text, IntWritable>{
6 public void reduce(Text tvname, Iterable<IntWritable> count, Context context)
7 throws IOException, InterruptedException{
8 int sum = 0;
9 for(IntWritable value: count){
10 sum+= value.get();
11 }
12 context.write(tvname,new IntWritable(sum));
13 }
14 }

```

Execution of Jar File:


```

[acadgild@localhost ~]$ hadoop jar Desktop/Jyoti/Assignment4 Task2.jar UnitCount /television.txt /output4_2
18/07/25 21:58:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/07/25 21:58:09 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/25 21:58:11 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface to execute your application with ToolRunner to remedy this.
18/07/25 21:58:11 INFO input.FileInputFormat: Total input paths to process : 1
18/07/25 21:58:11 INFO mapreduce.JobSubmitter: number of splits:1
18/07/25 21:58:12 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532534151354_0002
18/07/25 21:58:12 INFO impl.YarnClientImpl: Submitted application application_1532534151354_0002
18/07/25 21:58:12 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532534151354_0002/
18/07/25 21:58:12 INFO mapreduce.Job: Running job: job_1532534151354_0002
18/07/25 21:58:32 INFO mapreduce.Job: Job job_1532534151354_0002 running in uber mode : false
18/07/25 21:58:32 INFO mapreduce.Job: map 0% reduce 0%
18/07/25 21:58:46 INFO mapreduce.Job: map 100% reduce 0%
18/07/25 21:59:11 INFO mapreduce.Job: map 100% reduce 50%
18/07/25 21:59:14 INFO mapreduce.Job: map 100% reduce 100%
18/07/25 21:59:14 INFO mapreduce.Job: Job job_1532534151354_0002 completed successfully
18/07/25 21:59:15 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=210
    FILE: Number of bytes written=324328
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=834
    HDFS: Number of bytes written=38
    HDFS: Number of read operations=9
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=4
  Job Counters
    Killed reduce tasks=1
    Launched map tasks=1
    Launched reduce tasks=2
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=9128

```

```

    Total time spent by all reduces in occupied slots (ms)=48951
    Total time spent by all map tasks (ms)=9128
    Total time spent by all reduce tasks (ms)=48951
    Total vcore-milliseconds taken by all map tasks=9128
    Total vcore-milliseconds taken by all reduce tasks=48951
    Total megabyte-milliseconds taken by all map tasks=9347072
    Total megabyte-milliseconds taken by all reduce tasks=50125824
  Map-Reduce Framework
    Map input records=18
    Map output records=16
    Map output bytes=166
    Map output materialized bytes=210
    Input split bytes=101
    Combine input records=0
    Combine output records=0
    Reduce input groups=5
    Reduce shuffle bytes=210
    Reduce input records=16
    Reduce output records=5
    Spilled Records=32
    Shuffled Maps =2
    Failed Shuffles=0
    Merged Map outputs=2
    GC time elapsed (ms)=698
    CPU time spent (ms)=5180
    Physical memory (bytes) snapshot=398553088
    Virtual memory (bytes) snapshot=6186332160
    Total committed heap usage (bytes)=202575872
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=733
  File Output Format Counters
    Bytes Written=38
You have new mail in /var/spool/mail/acadgild

```

Output2:

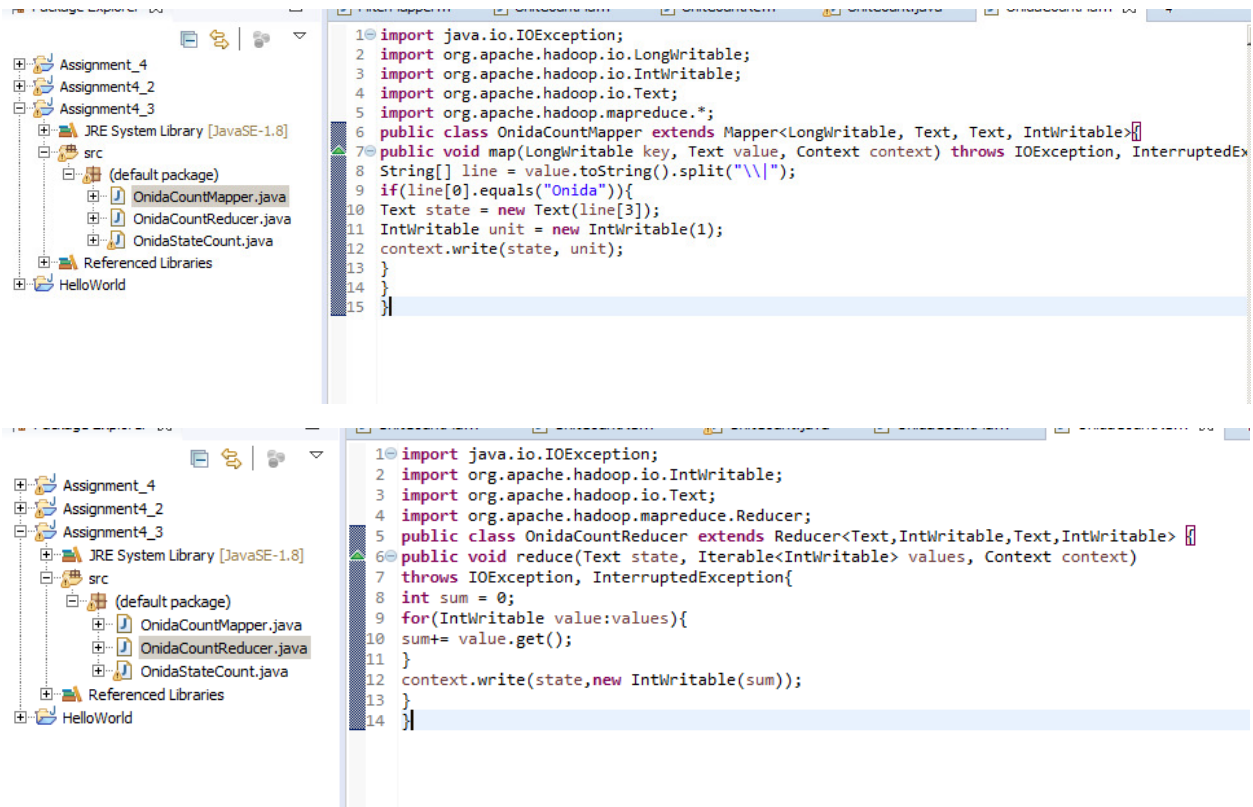
```
[acadgild@localhost ~]$ hadoop fs -cat /output4_2/part-r-00000
18/07/25 22:03:04 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform
Onida 3
Zen 2
[acadgild@localhost ~]$ hadoop fs -ls /output4_2
18/07/25 22:03:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform
Found 3 items
-rw-r--r-- 1 acadgild supergroup 0 2018-07-25 21:59 /output4_2/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 14 2018-07-25 21:59 /output4_2/part-r-00000
-rw-r--r-- 1 acadgild supergroup 24 2018-07-25 21:59 /output4_2/part-r-00001
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /output4_2/part-r-00001
18/07/25 22:03:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform
Akai 1
Lava 3
Samsung 7
[acadgild@localhost ~]$
```

Task 3:

Map Reduce program to calculate total units sold in each state for Onida Company.

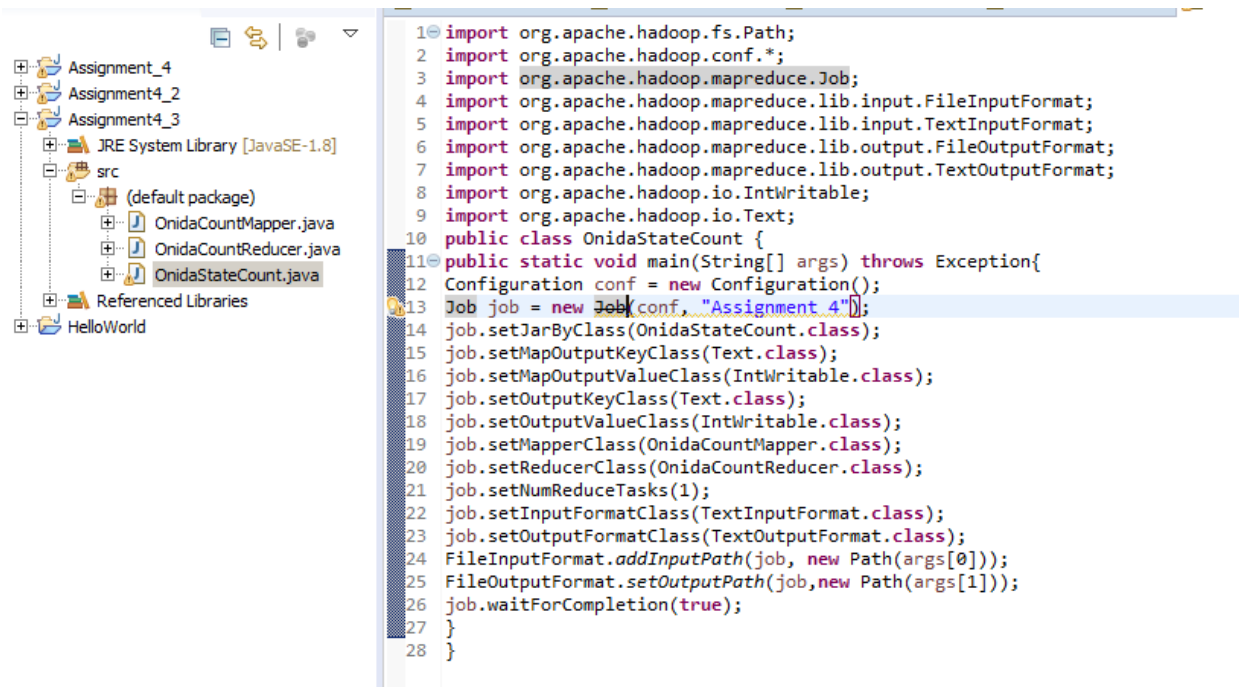
Solution:

Code:



```
1 import java.io.IOException;
2 import org.apache.hadoop.io.LongWritable;
3 import org.apache.hadoop.io.IntWritable;
4 import org.apache.hadoop.io.Text;
5 import org.apache.hadoop.mapreduce.*;
6 public class OnidaCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
7     public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
8         String[] line = value.toString().split("\\\\");
9         if(line[0].equals("Onida")){
10             Text state = new Text(line[3]);
11             IntWritable unit = new IntWritable(1);
12             context.write(state, unit);
13         }
14     }
15 }
```

```
1 import java.io.IOException;
2 import org.apache.hadoop.io.IntWritable;
3 import org.apache.hadoop.io.Text;
4 import org.apache.hadoop.mapreduce.Reducer;
5 public class OnidaCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
6     public void reduce(Text state, Iterable<IntWritable> values, Context context)
7         throws IOException, InterruptedException {
8         int sum = 0;
9         for(IntWritable value: values){
10             sum += value.get();
11         }
12         context.write(state, new IntWritable(sum));
13     }
14 }
```



Execution of Jar File:

```

[acadgild@localhost ~]$ hadoop jar Desktop/Jyoti/Assignment4 Task3.jar OnidaStateCount /television.txt /output4.3
18/07/25 22:17:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/07/25 22:17:17 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/25 22:17:19 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/07/25 22:17:19 INFO input.FileInputFormat: Total input paths to process : 1
18/07/25 22:17:20 INFO mapreduce.JobSubmitter: number of splits:1
18/07/25 22:17:20 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1532534151354_0003
18/07/25 22:17:22 INFO impl.YarnClientImpl: Submitted application application_1532534151354_0003
18/07/25 22:17:23 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1532534151354_0003/
18/07/25 22:17:23 INFO mapreduce.Job: Running job: job_1532534151354_0003
18/07/25 22:17:42 INFO mapreduce.Job: Job job_1532534151354_0003 running in uber mode : false
18/07/25 22:17:42 INFO mapreduce.Job: map 0% reduce 0%
18/07/25 22:17:57 INFO mapreduce.Job: map 100% reduce 0%
18/07/25 22:18:13 INFO mapreduce.Job: map 100% reduce 100%
18/07/25 22:18:13 INFO mapreduce.Job: Job job_1532534151354_0003 completed successfully
18/07/25 22:18:13 INFO mapreduce.Job: Counters: 49

File System Counters
  FILE: Number of bytes read=79
  FILE: Number of bytes written=216103
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=834
  HDFS: Number of bytes written=25
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=11809
  Total time spent by all reduces in occupied slots (ms)=12864
  Total time spent by all map tasks (ms)=11809
  Total time spent by all reduce tasks (ms)=12864
  Total vcore-milliseconds taken by all map tasks=11809

```

```

Total time spent by all map tasks (ms)=11809
Total time spent by all reduce tasks (ms)=12864
Total vcore-milliseconds taken by all map tasks=11809
Total vcore-milliseconds taken by all reduce tasks=12864
Total megabyte-milliseconds taken by all map tasks=12092416
Total megabyte-milliseconds taken by all reduce tasks=13172736
Map-Reduce Framework
  Map input records=18
  Map output records=4
  Map output bytes=65
  Map output materialized bytes=79
  Input split bytes=101
  Combine input records=0
  Combine output records=0
  Reduce input groups=2
  Reduce shuffle bytes=79
  Reduce input records=4
  Reduce output records=2
  Spilled Records=8
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=264
  CPU time spent (ms)=4750
  Physical memory (bytes) snapshot=297676800
  Virtual memory (bytes) snapshot=4117905408
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=733
File Output Format Counters
  Bytes Written=25
You have new mail in /var/spool/mail/acadgild

```

Output3:

```

[acadgild@localhost ~]$ hadoop fs -cat /output4_3/part-r-00000
18/07/25 22:19:18 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin
pplicable
Kerala 1
Uttar Pradesh 3
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```